

Cyber Assured Systems Engineering at Scale

Journal:	<i>IEEE Security & Privacy</i>
Manuscript ID	SPSI-2021-11-0255
Manuscript Type:	SPSI: May/June 2022: Formal Methods at Scale
Date Submitted by the Author:	30-Nov-2021
Complete List of Authors:	Cofer, Darren; Collins Aerospace Amundson, Isaac; Collins Aerospace Babar, Junaid; Collins Aerospace Hardin, David; Collins Aerospace Slind, Konrad; Collins Aerospace Alexander, Perry; University of Kansas School of Engineering Hatcliff, John; Kansas State University Robby, Robby; Kansas State University Klein, Gerwin; University of New South Wales Lewis, Corey; University of New South Wales Mercer, Eric; Brigham Young University, Computer Science Shackleton, John; Adventium Labs
Keywords:	D.2.4.d Formal methods < D.2.4 Software/Program Verification < D.2 Software Engineering < D Software/Software Engineering, D.4.6.g Verification < D.4.6 Security and Privacy Protection < D.4 Operating Systems < D Software/Software Engineering, D.3.2.h Development tools < D.3.2 Language Classifications < D.3 Programming Languages < D Software/Software Engineering, D.3.4.a Code generation < D.3.4 Processors < D.3 Programming Languages < D Software/Software Engineering

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60

Cyber Assured Systems Engineering at Scale

Darren Cofer, Isaac Amundson, Junaid Babar, David Hardin, Konrad Slind

Collins Aerospace

Perry Alexander

University of Kansas

John Hatcliff, Robby

Kansas State University

Gerwin Klein

Proofcraft and University of New South Wales, Sydney

Corey Lewis

University of New South Wales, Sydney

Eric Mercer

Brigham Young University

John Shackleton

Adventium Labs

Abstract—Formal methods tools that provide mathematical proof of system properties have improved dramatically in their power and capabilities. Our team has developed a model-based systems engineering environment that integrates formal methods at all levels of system design. Our methodology and tools enable systems engineers to address cybersecurity concerns early in the development of complex high-assurance systems.

■ **AEROSPACE SYSTEMS ENGINEERS** are currently given few development tools to understand and mitigate potential cybersecurity vulnerabilities. Typically, they must rely on process-oriented checklists and guidelines. Cyber vulnerabilities are often discovered during penetration testing late in the development process. Worse yet, they may be discovered only after the product has been fielded, necessitating extremely expensive and time-consuming remediation. This is not a sustainable development model.

Fortunately, formal methods tools have advanced to the point that they can be used to address cybersecurity and cyber-resiliency design challenges on real high-assurance systems at in-

dustrial scale, and do so much earlier in the development cycle. Our application domain is avionics and aerospace systems in general. This domain features large, real-time cyber-physical systems with the added complexities of performing safety-critical tasks as well as being exposed to a wide variety of cyber threats. Furthermore, aerospace systems are subject to intense regulatory scrutiny due to the certification requirements of this domain.

In previous work on the High-Assurance Cyber Military Systems (HACMS) project [1] we demonstrated that formal methods could be used to dramatically improve the cyber-resiliency of real aircraft, including an unmanned military helicopter. Our current work is focused on automating

the capabilities that we prototyped in the HACMS project and extending the reach and scale of the formal methods design and verification approach.

To this end, we have developed a model-based systems engineering (MBSE) environment that allows engineers to address a range of properties and manage system complexity through compositional analysis, integrating formal methods at all levels of the design process. MBSE processes utilize models as the primary vehicle for communication among the parties tasked with designing the system and as the primary design artifacts for requirements, verification, and code generation.

Our tools are based on the Architecture Analysis and Design Language (AADL) and extend the Open Source AADL Tool Environment (OS-ATE) [2]. The tools are specifically designed to bridge the gap between a user-level modeling language accessible to systems engineers and the highly specialized, formally verified code that implements the operating system (OS) kernel and other high-assurance components.

By using these tools to build real avionics systems, we show that current formal methods tools are practical, effective, and scalable to significant high-assurance applications in the aerospace industry.

INNOVATIONS

As part of the Cyber Assured Systems Engineering (CASE) project, our team has developed a MBSE tool environment that integrates design, verification, and code generation activities, enabling systems engineers to design-in cyber-resiliency for complex cyber-physical systems. The *BriefCASE* tools capture our vision for how formal methods can be applied throughout the design and build process to create high-assurance cyber-resilient systems.

A fundamental aspect of our approach is the use of architecture models to provide a framework for analyzing system behavior and organizing the assurance evidence produced. AADL allows engineers to describe the important elements of distributed, real-time, embedded systems (processors, memory, buses, processes, threads, and data interconnections) with sufficiently rigorous semantics that can support formal reasoning.

Proofs about models are meaningless unless there is some way to ensure that the implementa-

tion retains the properties of the model. The seL4 microkernel [3], used in both HACMS and CASE, is formally verified from its high-level security properties down to its binary implementation. By targeting seL4 we ensure that system components cannot interact in unintended ways and the *data flows* in the architecture model are enforced in the final product.

The main innovations of the BriefCASE tools and methodology are:

- 1) We provide automated architectural design patterns to address cyber-resiliency requirements, including synthesis of high-assurance components from formal specifications.
- 2) Our MBSE environment can target different operating systems including the seL4 microkernel, making its formal security guarantees easily accessible to developers. This ensures that the implementation produced is faithful to the modeled system.
- 3) Our approach is based upon co-evolution of system design and assurance artifacts, so that design changes automatically update the associated certification evidence. An *assurance case* is a structured argument, supported by evidence, intended to justify that a system is acceptably assured relative to a concern (such as safety or security) in the intended operating environment. An assurance case is embedded in the architecture model to capture and document the design decisions along with associated rationale.
- 4) Formal methods are integrated throughout the workflow, including requirements capture, component synthesis, verification, code generation, and the seL4 microkernel itself.

EXAMPLE

The example in Fig. 1 shows an AADL model of an unmanned air vehicle (UAV) for surveillance that was built with our BriefCASE tools. We will use this example to explain how the tools work together to implement the system and ensure its cyber-resiliency.

The system includes a ground station computer and the aircraft, consisting of a mission

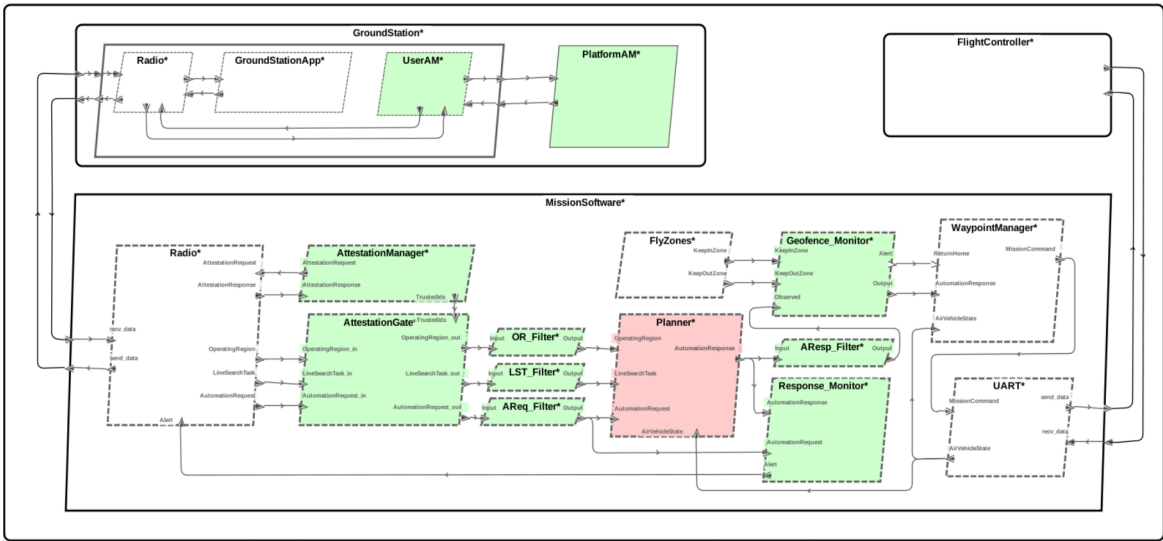


Figure 1. Cyber-resilient software architecture for UAV surveillance system.

computer and a flight control computer. The baseline (unhardened) mission computer included the four software components colored white: the radio for communication with the ground station (*Radio*), the mission planning service (*Planner*, provided as legacy software running on Linux), flight plan waypoint segmentation (*WaypointManager*), and a serial interface to communicate with the flight control computer (*UART*).

Cyber-threat analysis tools are used to analyze the unhardened functional model of the system, identifying the ground station and the mission planning service as primary sources of cyber attacks. Seven new cyber requirements are introduced that address ground station trust, message integrity, and run-time behavior vulnerabilities. The existing behavioral contracts in the unhardened system are strengthened to reflect these new requirements.

Design engineers use BriefCASE to transform the baseline system model to that shown in Fig. 1. Automated model transformations address the cyber requirements by inserting new high-assurance components (shown in green) and targeting the seL4 kernel to enforce separation between components. The *AttestationManager* establishes the trustworthiness of ground stations while the *AttestationGate* only passes messages from trusted sources. The three filters (*OR_Filter*, *LST_Filter*, and *AReq_Filter*) only

pass well-formed messages received from the *Radio*. Another filter (*AResp_Filter*) on the output of the *Planner* ensures that only well-formed flight plans are sent to the *WaypointManager*. Two run-time monitors alert the system to suspicious behaviors from the *Planner* such as flight plans that enter *keep-out* zones or leave *keep-in* zones (*Geofence_Monitor*), or unresponsiveness (*Response_Monitor*). The interface behavior of these high-assurance components, with the exception of the *AttestationManager*, is specified with assume-guarantee contracts (e.g., a filter makes no assumptions on input and only passes inputs that are syntactically well-formed).

The model is further transformed to move the mission planning service into a Linux virtual machine hosted on the seL4 microkernel. This permits us to run the legacy mission planner code without modification or porting and isolate any unintended behaviors. The target platform requires a static real-time schedule that is provided in the model. A transformation on the assume-guarantee contracts incorporates that schedule into the model for verification of the cyber requirements.

BRIEFCASE WORK FLOW

The BriefCASE environment provides systems engineers with a workflow and tool support for developing products with inherent cyber-

resiliency. In this section we provide an overview of the design, analysis, and code generation tools and how they can be used to implement high-assurance systems.

The workflow starts with the development of a baseline AADL model of the system architecture focusing on the desired functionality. This model can be analyzed using any of the existing AADL tools (e.g., resource usage, information flow, latency) to determine whether it is acceptable. BriefCASE integrates additional tools that analyze the architecture model for cybersecurity vulnerabilities and generate requirements that, when addressed, will mitigate those vulnerabilities. These requirements are imported into the model and may be addressed using a collection of automated model transforms. As requirements are addressed in the design, an assurance case is updated with corresponding evidence, computed directly from the model or by supporting analysis tools. Code implementing new high-assurance components as well as communication and execution infrastructure is generated from the model along with associated assurance evidence.

The following sections describe each step of the workflow in more detail.

Requirements

BriefCASE provides access to two analysis tools (GearCASE [4] and DCRYPPS [5]) that can examine AADL models to detect potential cyber vulnerabilities and suggest requirements for mitigation. Systems engineers are presented with a requirements management interface (top pane in Figure 2) for viewing the generated requirements and importing them into the model so they can be addressed. The interface enables engineers to select the requirements they wish to import and assign them unique identifiers, or omit them with rationale. A document listing the omitted requirements and rationale is maintained and may be a required development artifact for some certification domains.

Some requirements can also be formalized as assume-guarantee contracts added to the AADL model, enabling formal verification. Such a requirement will be imported into the model with with an associated formal contract.

A BriefCASE project contains a repository for requirements. Imported requirements are repre-

sented as assurance case goals to be satisfied. For example, one of the requirements for well-formed messages (selected in Fig. 2) is imported as the following goal:

```
goal Req_WellFormed_OperatingRegion(comp_context : component) <=
  ** "UxAS component shall only receive well-formed messages" **
context Generated_On : "January 29, 2021";
context Req_Component : "MC::MissionComputer.Impl.SW.UxAS";
undeveloped
```

The goal is marked *undeveloped* initially. Evidential statements are added to the goal as the design is updated to address this requirement.

Cyber Transforms

To address the new cyber-resiliency requirement, the architecture will need to be transformed in such a way as to harden the design against the vulnerability. BriefCASE provides a library of model transformations for addressing common cyber vulnerabilities. Currently, the following transformations are supported:

- Filter – Blocks messages that do not conform to the given specification
- Monitor – Detects violations of a given runtime condition and generates an alert
- Switch – Used with a Monitor to block messages when an alert is generated (also referred to as a *gate*)
- Attestation – Performs a measurement on non-local software to assess its trustworthiness
- Virtualization – Isolates software component(s) in a virtual machine
- Proxy – Inserts a pair of components to allow inspection of HTTPS message payloads
- seL4 – Transforms the model to comply with seL4 component properties

The transformations are automated by the BriefCASE tool, resulting in a hardened model that is correct-by-construction. For example, the requirement that a component only receives well-formed messages can be satisfied by the insertion of a high-assurance filter. A BriefCASE transform wizard helps to configure the filter component properties, including the filter behavioral specification, which is represented as an assume-guarantee contract.

BriefCASE then inserts a new filter component into the model, sets the component properties, and establishes the appropriate connections to source and destination components. The filter behavioral contract is also added to the model,

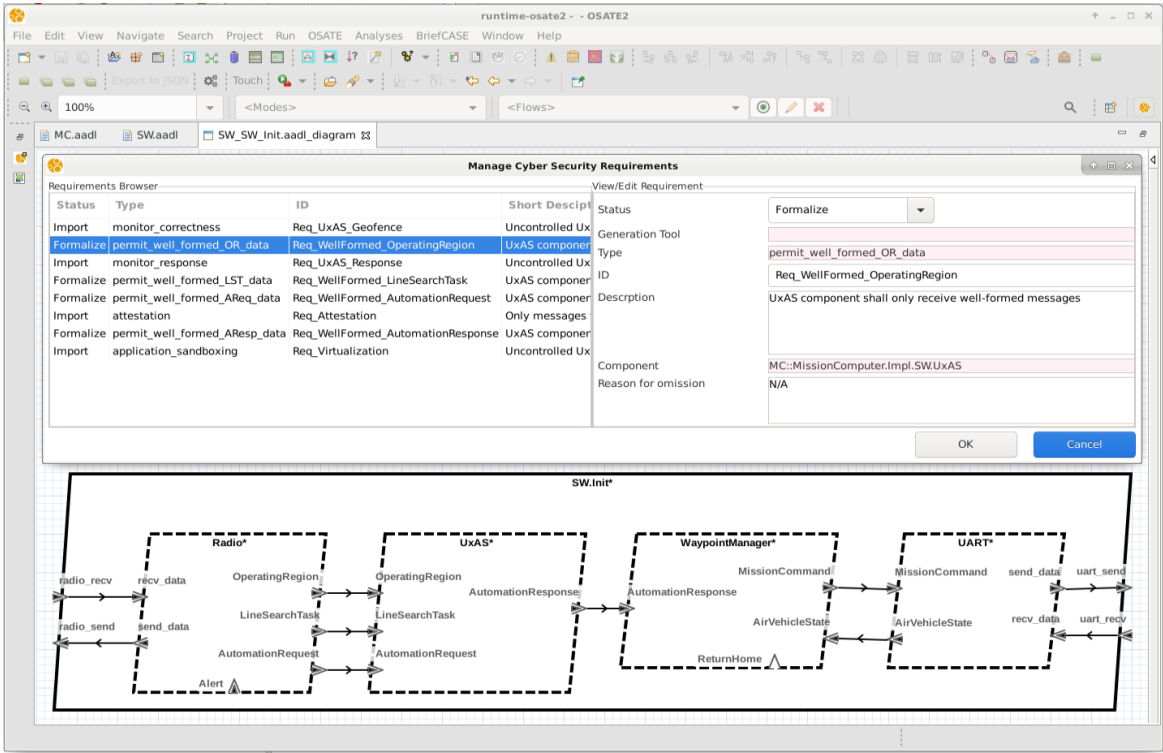


Figure 2. BriefCASE requirements management interface.

enabling formal analysis of the model as well as providing the behavioral specification for a provably correct synthesis of the filter component implementation.

The transformation also updates the assurance case with new evidential statements indicating that the associated goal has been satisfied, including the strategy used and links to context and associated evidence.

Compositional Analysis

The Assume Guarantee Reasoning Environment (AGREE) is a compositional, assume-guarantee-style model checker for AADL models [6]. AGREE attempts to prove properties about one layer of the architecture using properties allocated to subcomponents. The composition is performed in terms of assumptions and guarantees that are provided for each component. Assumptions describe the expectations the component has on the environment, while guarantees describe bounds on the behavior of the component.

AGREE uses k-induction as the underlying algorithm for model checking. AADL models

and AGREE contracts are first translated into the Lustre language, including verification conditions for consistency and correctness of the contracts. The model checker then attempts to find any model execution traces that would violate these conditions using one of several Satisfiability Modulo Theories (SMT) solvers. If the model checker covers all reachable states in the model without finding a violation then the properties are proved.

Once the system architecture has been modeled using AADL and the component behavior is specified using AGREE’s assume-guarantee contracts, we use the AGREE model checker to verify the consistency of these contracts.

- 1) Component interfaces – The output guarantees of each component must be strong enough to satisfy the input assumptions of downstream components.
- 2) Correctness of implementations – The input assumptions of a system along with the output guarantees of its sub-components must be strong enough to satisfy its output

guarantees.

For example, in Fig. 1, the input assumptions of the Waypoint Manager must be weaker than the output guarantees of the Geofence monitor and the output guarantees of the Mission Software must be inferrable from its input assumptions combined with the output guarantees of its components.

This hierarchical strategy for reasoning about contracts, or *compositional analysis*, reduces the computational complexity of model checking by breaking down the larger problem into more manageable fragments. Verification of a system does not directly depend on the implementations of its components (or their sub-components), but only on their contracts. In Fig. 1, each of the Mission Software components either contains subcomponents or has a corresponding source code implementation of its contract. Rather than reasoning monolithically about the interaction of all of these lower-level implementations, the compositional approach allows us to reason about each layer of the model separately.

High Assurance Component Synthesis

The correctness of the high-assurance components inserted by BriefCASE transformations means that each such component must meet its AGREE contract. This obligation is addressed by *formal synthesis*, using the Semantic Properties of Language and Automata Theory (SPLAT) tool. Given a sufficiently detailed formal specification of component behavior, SPLAT generates code, as well as proofs showing that the generated code is correctly compiled and meets its specification [7].

The formal languages that filters, gates, and monitors are generated from include regular expressions, contiguity types. For each of these languages, we have infrastructure that (a) translates formal specifications to code and (b) proves the correctness of the translation using the HOL4 theorem proving system. The generated code is *CakeML*, a dialect of Standard ML possessing a fully verified compiler [8].

Regular expressions and contiguity types are formalisms well-suited for expressing classes of constraints on the data passed in messages. Thus, they provide a useful basis for expressing filters. Regular expressions compile to efficient

finite state automata, generating a correctness proof along the way. We have also synthesized regular expressions to hardware using this proof-producing technique. Contiguity types can capture more complex data formats, including variable-length arrays and union structures which are used in the Mission Planner found in our UAV example. We deploy a verified contiguity-type based parser generator to decode and check that the serialized message data is well-formed relative to its specification.

Gates and monitors can require arbitrarily complex computations for their work; hence we also support code generation from Lustre. We have formalized the semantics of Lustre and are using it as the basis for a formal translation to CakeML.

Since code generated from SPLAT is deployed as a scheduled thread in a real-time environment, there are two aspects to consider: (1) correctness of a ‘one shot’ thread execution (discussed above), and (2) correctness of the perpetual re-execution of the thread. Modeling the latter and proving relevant properties is discussed in [9].

Remote Attestation

Semantic remote attestation is a technique for establishing trust in software running on a non-local computer. An appraiser requests an attestation from a target, receives evidence in response to the request, and appraises the evidence to determine trust [10]. Because remote attestation does not require modification of its measurement target, we utilize it to establish trust in legacy software that cannot otherwise be verified. We construct a verified remote attestation infrastructure around the legacy target that generates run-time and boot-time evidence.

In the UAV example, we wish to ensure that the aircraft will only accept commands from trustworthy (non-compromised) ground stations. Our approach adds three attestation managers to the seL4-based architecture in Fig. 1. Each attestation manager executes protocols specified by Copland [11] phrases. Copland is a formal specification language designed for writing attestation protocols that are both verifiable and executable. The attestation managers themselves are written and verified using CakeML and Coq.

The *AttestationManager* on the Mission Com-

puter makes attestation requests and appraises the results. A protocol request and nonce are sent to the remote target (in this case, the Ground Station), evidence is returned, and results appraised to determine trust. If the appraisal is successful, the identifier for the target is added to a list of trusted computers whose messages will be allowed to pass through the *AttestationGate*.

Two other attestation managers must be added on the Ground Station. The Ground Station software has been modified to run in a Linux virtual machine hosted on seL4. The UserAM runs as a Linux process on this virtual machine. Its responsibilities include responding to attestation requests from the UAV, measuring the Ground Station application software, and requesting attestations from the Ground Station platform (the Linux kernel). When the UserAM receives an attestation request it responds by executing an attestation protocol that measures the application, requests measurements from the PlatformAM, signs the result, and generates a nonce from the request. The resulting evidence package is returned to the requesting appraiser on the UAV.

Because the UserAM runs as a Linux process, it cannot be trusted *a priori*. A PlatformAM is introduced to perform an attestation of the Linux kernel (and the UserAM). The PlatformAM runs as a seL4 component outside of the Linux virtual machine separated from the Linux kernel. seL4's guaranteed separation properties provide assurance that the PlatformAM cannot be compromised by other platform software. The PlatformAM only responds to requests from the UserAM and similarly runs a protocol that produces signed results.

Using the remote attestation system to harden a platform involves writing application-specific attestation and appraisal components. New measurers are constructed for the UserAM that measure the running application along with corresponding appraisal code. The PlatformAM and attestation architecture remain the same across applications. Thus, the overhead required for the implementation is minimized.

Information Flow Analysis

As systems become more complex, composition of components and systems presents safety and security challenges that span many

sub-systems. Subsystems and components may originate from different organizations, many of which may not have a complete understanding of system information flows and potential impacts of security attacks. Thus, it is essential to have trustworthy methods of developing common understanding of dependencies in the system and the respective responsibilities of the involved vendors.

The Awas [12] AADL information flow analyzer and visualizer has been applied to enable developers and auditors to understand, reason, explore, and visualize system dependencies and information flow at scale across components and sub-systems. Awas processes the AADL system architecture model, specifically its inter-component connection descriptions and intra-component flow specifications, to provide formal system-wide impact and flow analyses. Such flows include component data/control flows, security-oriented information flows, and fault/error propagation specified using the AADL Error Modeling Annex (EMv2). Awas also provides a user-friendly Domain Specific Language (DSL) to query, check, and visualize custom safety/security system properties.

Fig. 3 illustrates an example of Awas information flow reachability analysis. In our UAV example, Awas can compute and visualize how the information in Ground Station messages flows through the system as well as the components or ports that may directly or indirectly consume data derived from that information. Awas supports a number of forms of forward and backward interactive information queries. Using the Awas script-based query language, one can specify and check more complex properties, e.g., that information must flow through specific ports or components. These end-to-end flow specifications are often useful for supporting verification of the effectiveness of cyber-resiliency components. An Awas specification can state that information from untrusted components such as the Ground Station always flow through the Attestation Gate and filter components before reaching the Flight Control or other components that make critical decisions about the flight or mission of the UAV.

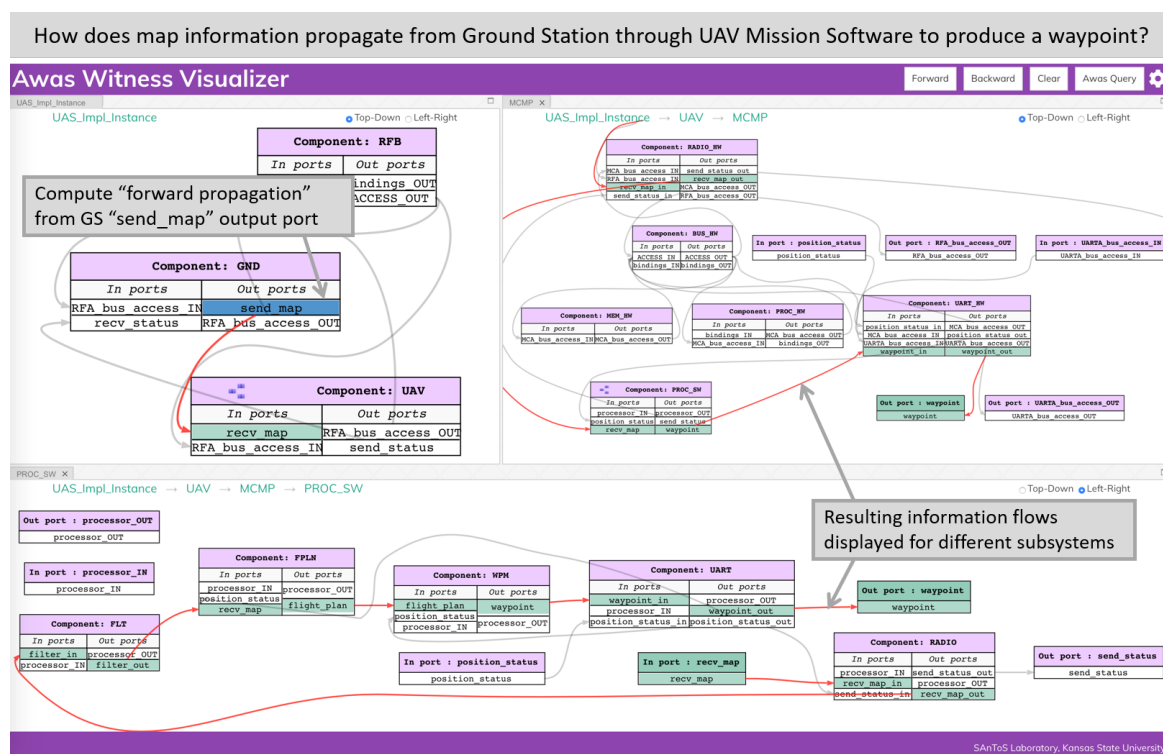


Figure 3. Awas Information Flow Analysis.

Real-Time Scheduling

In mixed-criticality mission systems, execution threads often have strict confidentiality, integrity, and availability requirements. Under certain conditions, timing channels can be manipulated to violate these security requirements. For example, temporal interference between timing channels can reduce availability of time-critical functionality and weaken the integrity of controls by inducing selective jitter. As a result, a compromised component could dominate the processor, preventing other components from completing their tasks, placing the security of the entire platform in jeopardy. Simple scheduling approaches that attempt to use priority schemes to mitigate this impact only protect the highest priority threads and leave the other threads vulnerable.

Temporal isolation is a more secure technique to restrict timing channels and reduce temporal interference between software threads executing on the same platform. BriefCASE achieves temporal isolation by leveraging prior work from safety and security-critical disciplines, such as avionics, where temporal isolation in real-time

scheduling has been deployed for decades. We implemented a static cyclic scheduling approach using the seL4 domain scheduler, where a fixed schedule defines an ordered sequence of static execution slots. Each slot has a duration and a partition identifier. seL4 ensures that the temporal domain, and any threads running in it, will not exceed their defined time allotments. Thus, each thread executes within its own well-defined timing constraints without interference from other threads. System designers are responsible for the creation of a valid static cyclic schedule that satisfies all the application’s timing requirements.

BriefCASE generates a start-of-frame synchronization signal for each thread using a special thread called the Pacer, which sends periodic signals to each thread. Each application thread blocks until it receives its signal from the Pacer. The application thread runs to completion, and then blocks again on the Pacer signal for the next iteration. Each thread subsequently executes exactly once during its statically scheduled time slice.

The new seL4 Mixed-Criticality Systems

(MCS) variant provides additional capability that can support temporally isolated real-time systems. As part of BriefCASE we developed a proof-of-concept static cyclic scheduler for MCS. It includes a start-of-frame signal, which eliminates the need for the Pacer component. It also includes kernel level support for flexible dynamic scheduling that satisfies some real-time properties, such as period and execution time.

Infrastructure Code Generation

The High Assurance Modeling and Rapid engineering for embedded systems tool (HAMR) [13] is a multi-platform, multi-language AADL code generation framework. In the CASE project, HAMR is primarily used to generate code for deployment for the seL4 microkernel, but system and component prototyping is also supported utilizing HAMR’s code generation capability for the Java Virtual Machine (JVM) and Linux. One of the primary objectives is to support system builds that leverage seL4 micro-kernel partitioning and information flow guarantees to achieve the AADL-specified component separation and inter-component communication needed for cyber-resiliency.

For each AADL thread component, HAMR generates a thread code skeleton and application programming interfaces (APIs) for communicating over the ports declared on the component. For components that are implemented manually, the developer fills out the thread skeleton with application code, calling the port APIs, libraries, and component-specific developer-implemented support code as needed. HAMR supports coding component application logic in either C, Slang [14] (a high-assurance subset of Scala that can be translated to C), or CakeML. Slang-implemented components can interface with Scala or Java libraries and execute on the JVM. Pure Slang components can be translated to C and then deployed on HAMR-built Linux or seL4 systems. On CASE, Slang has been used for prototyping components, but primary development of directly implemented components has used C. For high-assurance components specified using SPLAT, HAMR automatically integrates code generated from SPLAT into the component binary code using CakeML’s Foreign Function Interface (FFI) mechanism.

HAMR generates component infrastructure and integration code implementing the semantics of AADL-compliant thread scheduling, thread dispatching, and port-based communication. For port communication, both shared memory communication (AADL data ports), buffered messaging (AADL event data ports), and buffered notification (AADL event ports) are supported. HAMR code generation is staged using a translation architecture that facilitates adding new backends for different target platforms. Almost all the infrastructure code is implemented in Slang, which can then be used for JVM deployments or translated to C for Linux or seL4 deployments. The Slang-based implementation of the AADL run-time framework can be viewed as a high-level reference implementation of AADL semantics. Automatic translation of this reference implementation to C on different platforms helps establish semantic consistency across those platforms. The AGREE contract language and SPLAT code generation framework have been designed to align with the AADL semantics reflected in this reference implementation.

To generate deployments on seL4, HAMR makes heavy use of the component architecture for microkernel-based embedded systems (CAMkES) code-generation framework. The CAMkES input language supports specification of seL4 partitioning and communication topology using component-oriented idioms. The CAMkES translator generates an seL4 *Capability Description Language (CapDL)* file that configures the seL4 kernel to support protected memory blocks and permission-based reading and writing of each block as indicated by CAMkES components and connections. To realize the memory separation specified by the AADL architecture description, HAMR generates CAMkES specifications that: (a) reflect the AADL model’s component and communication topology, and (b) include additional components and communication to support the AADL run-time infrastructure, in particular, thread scheduling. To enforce time partitioning, HAMR uses the seL4 domain scheduler to support static cyclic scheduling.

HAMR also provides automated support for configuring Linux-based virtual machines as components within the seL4 deployed system. In the example UAV system, this feature has been

used to sandbox untrusted legacy code for the Mission Planner component.

HAMR-generated seL4 systems can be executed on the QEMU emulator before deployment to a development board or production hardware. This prototyping capability significantly speeds up development iterations and enables the development of a sophisticated automated regression testing framework.

While we have not yet completed full formal verification of the HAMR Slang-based reference implementation and code generation pipeline, information flow preservation is a key property that illustrates how one might approach formal verification of other important semantic properties. Specifically, HAMR generates flow graphs reflecting the inter-component information flow at both the AADL architecture level and the seL4 CAMkES level. In addition to multiple visualizations and auditable traceability artifacts, HAMR generates SMT-based representations of the flow graphs and traceability relationships between them. Formalized theorems for information flow preservation can then be represented as SMT assertions. This enables SMT solvers to automatically prove the following properties for any HAMR-generated seL4 deployment:

- 1) All AADL flows are implemented – For every inter-component information flow present in the AADL model, that flow is provisioned in the generated seL4 kernel configuration.
- 2) No extraneous flows are implemented – For each inter-component flows provisioned in the seL4 kernel, that flow is represented as an inter-component connection in the AADL model.

These properties were chosen to focus on cyber-resiliency needs, but other key semantic properties of the AADL run-time can be incrementally formally verified, e.g., by applying the Slang Logika formal verification capability to the Slang reference implementation of the AADL run-time.

HAMR plays a key role in integrating formal methods at different levels of abstraction at multiple points throughout the development process, enabling those methods to be applied at-scale in the system development process. Using the formally verified seL4 microkernel as a

foundation, HAMR enables AADL to be used as a model-based development and systems engineering framework for seL4-based applications. HAMR provides a semantically-consistent multi-platform code generation process that enables: (a) formally verified components (e.g., generated from SPLAT) to be correctly integrated and deployed; and (b) formal specification and verification frameworks like AGREE to be used to reason about both component and system level properties. The HAMR code generation architecture is designed to support strong traceability and verification, as illustrated by the information flow correspondence properties. The compositional and staged nature of HAMR-based development supports scaling of formal approaches by enabling them to be included in a component-wise manner and at different levels of abstraction, while also integrating parts of the implementation that are assured using traditional methods. In addition, the strong partitioning of seL4 enables the controlled integration of untrusted components.

Secure Microkernel

The seL4 microkernel [3] is a lightweight, fast, and secure operating system kernel. Its implementation is fully formally verified, from high-level security properties down to the binary level. It was the first OS kernel with this degree of formal verification, and after more than a decade of further research and engineering is still not only the leading formally verified OS kernel, but also the fastest OS kernel on the Arm architecture.

Its formal verification makes seL4 the ideal basis for high-assurance systems. It is in itself a demonstration of the level of fidelity and scale formal verification can achieve [15]. seL4 supports multiple architectures (Arm, x86-64, RISC-V), provides deep security properties such as integrity, confidentiality and availability, and comes with formally verified user-level system initialization. As one of its multiple available OS personalities, seL4 also offers the user-level CAMkES component system that provably achieves isolation between statically specified components.

The formal proofs about seL4 and the corresponding user-level components measure over one million lines of proof script in total. They constitute one of the largest continually maintained formal proof artifacts in existence and

1 provide a rich target for new techniques in proof
2 engineering, proof repair, and automation for con-
3 structing new proofs about software as well as
4 maintaining existing large-scale proof artifacts.

5 While it is essential to build a high-assurance
6 system on a high-assurance OS kernel, this is
7 not the main feature seL4 provides for systems
8 engineering — a simpler real-time OS might be
9 formally verified, but would not be sufficient for
10 the engineering method described in this paper.
11 The true power of seL4 lies in its ability to scale
12 formal analysis and verification to the much larger
13 code bases that make up entire systems. It does so
14 by providing strong isolation between user-level
15 components [16].

16 This isolation means that components can
17 be analyzed separately from each other and be
18 composed safely — in this way seL4 provides
19 the foundation that the soundness of the highly
20 automated analysis tools such as AGREE depend
21 on. It makes it possible to run entire untrusted
22 virtual machines and securely monitor their be-
23 havior on the same hardware. It makes it possible
24 to provide filter and monitor components and
25 prove that these components cannot be tampered
26 with by the components they protect. And it
27 makes it possible to guarantee that the limited
28 communication channels that the analysis tools
29 assume to exist in the AADL model are the *only*
30 communication channels that are available to the
31 components in the system. The combination of
32 these enables automated high-level analysis with
33 high assurance.

34 Assurance Case

35 An important aspect of our work has been to
36 structure formalizations and proofs by following
37 the AADL model of the system. We have found
38 that in assuring the cyber-resiliency properties
39 of aircraft designs we need to integrate different
40 kinds of evidence with varying levels of formal-
41 ity. This has been our motivation for exploring
42 assurance case methods.

43 In previous work, we developed the *Resolute*
44 language and tool [17] as a way to help engi-
45 neers create an assurance argument describing the
46 steps taken during the design process to make
47 the system safe and secure. The Resolute syn-
48 tax supports construction of assurance cases that
49 comply with the Goal Structuring Notation (GSN)

v2 standard. Claims are expressed as *goals* and
strategies, and can contain attributes such as *con-*
text, *assumptions*, and *justification*. Claims can be
marked *undeveloped*, which Resolute interprets as
an unsupported claim, or with a *solution*, which is
an explicit assertion that the claim is supported.
Rather than being a separate document, a Res-
olute assurance case is part of the architecture
model and can refer to elements within the model.
Since it is not a static representation, it can ensure
that the assurance argument remains consistent
with the evolving design.

A partial assurance case for the hardened
UAV surveillance system generated by Resolute is
shown in Fig. 4. It illustrates assurance argument
subtrees for several of the high-assurance com-
ponents associated with cyber-resiliency require-
ments. There is also a large subtree associated
with HAMR code generation and traceability to
seL4 separation guarantees that is collapsed in
this view due to size constraints.

BriefCASE includes a library of Resolute
assurance strategies, or *patterns*, that align with
the CASE workflow. The patterns are instantiated
with context from the AADL model and spec-
ify the evidence required to support the cyber-
resiliency goals of the system. For example, the
`add_filter` strategy is automatically inserted
into the assurance case when the *Filter* trans-
formation is performed, and includes logical rules
that Resolute uses to determine whether the claim
of well-formed messages is supported by evi-
dence. The `add_filter` definition includes the
following sub-goals:

- `filter_exists` – The filter component is still present in the model and has not been altered or deleted by subsequent design changes.
- `filter_not_bypassed` – There is no alternate information flow in the model that would allow the filter to be bypassed and therefore not perform its function.
- `filter_implemented_correctly` – The filter has been implemented correctly to meet its AGREE specification.

The first two sub-goals are supported by evi-
dence obtained by examining the structure of
the model. If at a later time during development
the model is inadvertently altered in a way that
renders the transformation ineffective, Resolute

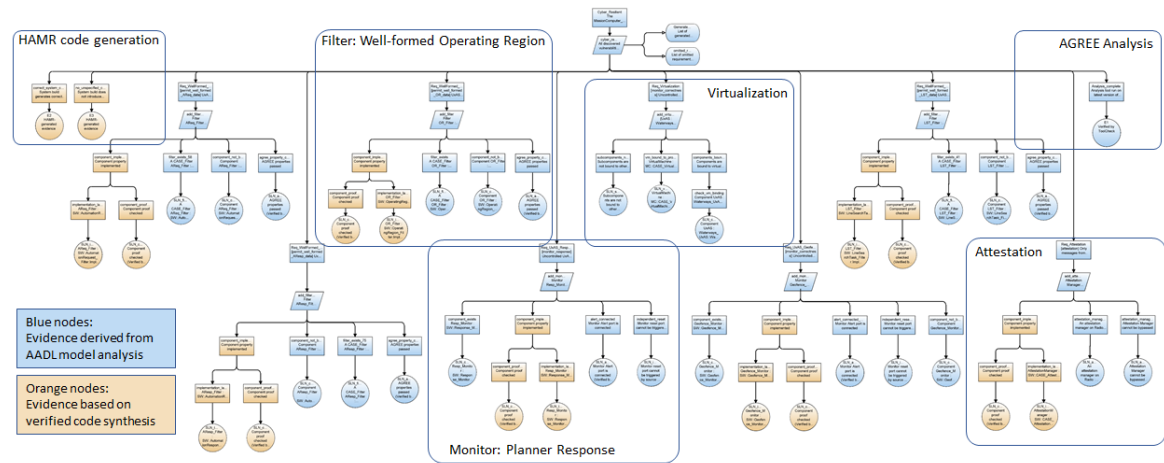


Figure 4. Resolute assurance case for hardened UAV system (partial).

will be unable to substantiate the evidential statements and will produce a failing assurance case. For the third subgoal, Resolute uses the existence of the SPLAT proof corresponding to the filter's AGREE specification as evidence that the component was implemented correctly.

AIRCRAFT APPLICATION

We have successfully demonstrated our BriefCASE methodology and tools to develop proof-of-concept enhancements for the CH-47F helicopter Common Avionics Architecture System (CAAS). As an integrated cockpit avionics suite developed by Collins Aerospace, CAAS serves as a prime example of modern air platform complexity with common avionics across a variety of defense platforms including the Army's Mission Enhanced Little Bird (MELB) and MH-60, the Navy CH-53K, and the Air Force KC-135. The CAAS application offered an opportunity to apply BriefCASE tools across a variety of mission systems, including legacy components, flight critical software, as well as new and evolving systems.

Most of the work for this demonstration was performed by a team of CAAS development engineers who had no previous experience with formal methods. Our primary goal was to enable CAAS product engineers to employ the CASE tools to create an operational mission scenario exhibiting cyber threats and mitigations that could be exercised using the facilities of the Collins CH-47F System Integration Laboratory.

The CH-47F demonstration system for CASE focused on integrating pilot and soldier wireless tablet computers for increased situational awareness and display of Automatic Dependent Surveillance-Broadcast (ADS-B) data regarding nearby air traffic (Fig. 5). BriefCASE tools were used to implement this networking enhancement while ensuring that no new cybersecurity vulnerabilities were introduced. A high-assurance gateway was added between the existing CAAS network and the new wireless network, including new components for monitoring messages to and from the wireless devices. Remote attestation was also added to ensure that any devices that attempt to join the wireless network are running trustworthy software. This also required configuring the seL4 microkernel to run on an existing CAAS processing module (the vision processing module, or VPM) that was repurposed to serve as the secure gateway.

The CAAS engineers first developed an AADL model of the CH-47F CAAS system. They added the enhanced capabilities for wireless access described above, resulting in a baseline architecture. The engineers then analyzed their baseline AADL architecture utilizing the GearCASE and DCRYPPS tools, resulting in a set of cyber requirements. They employed the BriefCASE AADL tools to add filters, monitors, attestation gates, and seL4 isolation to satisfy the cyber requirements. ADS-B anti-spoofing monitors were also developed by the CAAS team

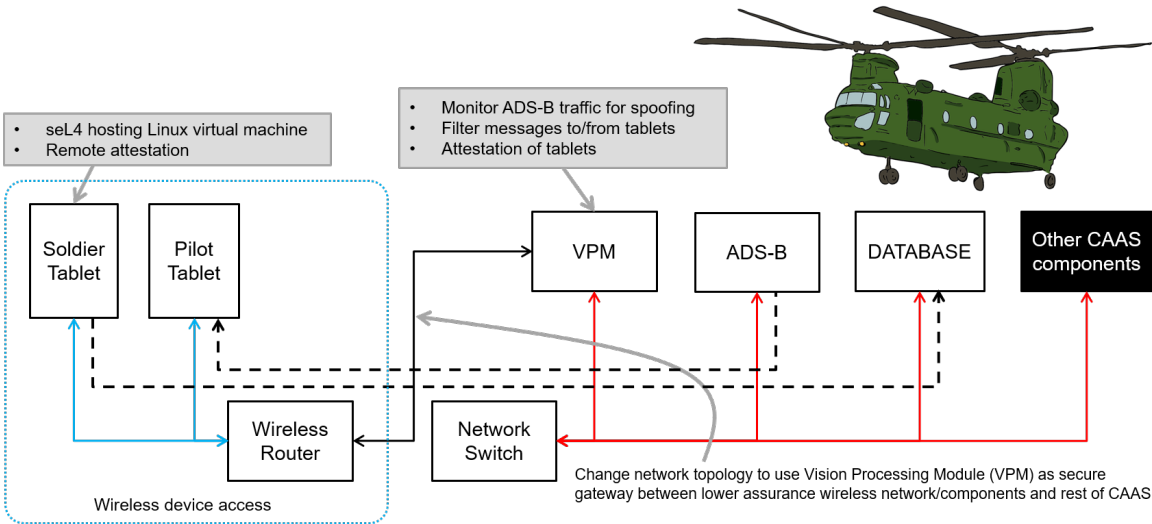


Figure 5. Cyber-resilient software architecture adding wireless device access to CH-47 avionics.

to detect inconsistencies in aircraft position and velocity trends, bad traffic identifiers, and other possible indicators of spoofed aircraft.

The specifications for the filters and monitors were developed by the CAAS engineers using the AGREE contract language, with assistance from our research team. The SPLAT tool was used to synthesize the monitors and filters from the AGREE specifications with high assurance, as described above. The HAMR tool was then used to generate infrastructure code for the overall system running on seL4.

The tablet operating environment was modified to run application software in a Linux virtual machine hosted on seL4. This was done so that the remote attestation component for platform measurement could be securely isolated from the (untrusted) application software.

An additional complication was the need to add proxy and network adapter components to allow the encrypted internet messages from devices on the untrusted wireless network to be examined by the secure gateway. Dedicated processing cores were available on the VPM for the provisioning of these “low-side” and “high-side” components separately from the other high-assurance software. The architecture model was transformed to add these components but their implementation code had to be manually written due to the complexity of dealing with off-the-

shelf networking technologies. Automating the synthesis of network adapter and proxy components is future work.

As initial users of the BriefCASE tools, the CAAS engineers encountered some limitations in the specification expressiveness and documentation, as might be expected for the first use of research tools by product area developers. This led to improvements in the tools to add or extend capabilities. Initial estimates of processing time required for the complex monitoring components turned out to be optimistic, requiring an optimization effort. Lack of detailed documentation and hardware instabilities hampered the tablet software development effort. But overall the BriefCASE tools were able to be productively used by Collins product engineers to produce the CH-47F CAAS demonstration system on time and within budget, providing our research team with valuable feedback on the strengths and weaknesses of the current BriefCASE tool environment.

CONCLUSION

We have produced a model-based systems engineering environment called BriefCASE, based on the Architecture Analysis and Design Language. The BriefCASE tools and methodology provide design, analysis, and code generation capabilities based on formal methods and targeted at high-assurance cyberphysical systems.

Key innovations include automated architectural design patterns for cyber-resiliency, co-evolution of system design and assurance artifacts (captured as an assurance case linked to the architecture model), synthesis of code for high-assurance components, and code generation targeting the formally verified seL4 microkernel.

We have demonstrated the effectiveness and scalability of these tools by using them to add new cyber-resilient features to a military helicopter avionics system. Their successful use by our product engineers provides evidence that formal methods can be incorporated into industrial projects.

All of the tools are open-source, with links and documentation available at <http://loonwerks.com/projects/case.html>. We hope that others will find value in this approach and extend the tools with new cyber transforms, expanded system analysis tools, and code generation for additional operating systems and computing platforms.

ACKNOWLEDGMENT

This work was funded by DARPA contract HR00111890001. The views, opinions and/or findings expressed are those of the authors and should not be interpreted as representing the official views or policies of the Department of Defense or the U.S. Government.

REFERENCES

1. D. D. Cofer, A. Gacek, J. Backes, M. W. Whalen, L. Pike, A. Foltzer, M. Podhradsky, G. Klein, I. Kuz, J. Andronick, G. Heiser, and D. Stuart, "A formal approach to constructing secure air vehicle software," *Computer*, vol. 51, no. 11, pp. 14–23, 2018.
2. P. H. Feiler and D. P. Gluch, *Model-Based Engineering with AADL: An Introduction to the SAE Architecture Analysis & Design Language*, 1st ed. Addison-Wesley Professional, 2012.
3. G. Klein, K. Elphinstone, G. Heiser, J. Andronick, D. Cock, P. Derrin, D. Elkaduwe, K. Engelhardt, R. Kolanski, M. Norrish, T. Sewell, H. Tuch, and S. Winwood, "seL4: formal verification of an OS kernel," in *Proceedings of the 22nd ACM Symposium on Operating Systems Principles 2009, SOSP 2009, Big Sky, Montana, USA, October 11-14, 2009*, J. N. Matthews and T. E. Anderson, Eds. ACM, 2009, pp. 207–220. [Online]. Available: <https://doi.org/10.1145/1629575.1629596>
4. T. Patten, D. Mitchell, and C. Call, "Cyber attack grammars for risk-cost analysis," in *Proceedings of the 15th International Conference on Cyber Warfare and Security*, Norfolk, VA, 2020.
5. R. Laddaga, P. Robertson, H. E. Shrobe, D. Cerys, P. Manghwani, and P. Meijer, "Deriving cyber-security requirements for cyber physical systems," *CoRR*, vol. abs/1901.01867, 2019. [Online]. Available: <http://arxiv.org/abs/1901.01867>
6. D. Cofer, A. Gacek, S. Miller, M. W. Whalen, B. LaValley, and L. Sha, "Compositional verification of architectural models," in *NASA Formal Methods*, A. E. Goodloe and S. Person, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 126–140.
7. E. Mercer, K. Slind, I. Amundson, D. Cofer, J. Babar, and D. Hardin, "Synthesizing verified components for cyber assured systems engineering," in *24th International Conference on Model-Driven Engineering Languages and Systems (MODELS 2021)*, October 2021.
8. R. Kumar, M. O. Myreen, M. Norrish, and S. Owens, "CakeML: a verified implementation of ML," in *The 41st Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '14, San Diego, CA, USA, January 20-21, 2014*, S. Jagannathan and P. Sewell, Eds. ACM, 2014, pp. 179–192. [Online]. Available: <https://doi.org/10.1145/2535838.2535841>
9. J. Å. Pohjola, H. Rostedt, and M. O. Myreen, "Characteristic formulae for liveness properties of non-terminating CakeML programs," in *10th International Conference on Interactive Theorem Proving, ITP 2019, September 9-12, 2019, Portland, OR, USA*, ser. LIPIcs, J. Harrison, J. O'Leary, and A. Tolmach, Eds., vol. 141. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2019, pp. 32:1–32:19. [Online]. Available: <https://doi.org/10.4230/LIPIcs.ITP.2019.32>
10. G. Coker, J. Guttman, P. Loscocco, A. Herzog, J. Millen, B. O'Hanlon, J. Ramsdell, A. Segall, J. Sheehy, and B. Sniffen, "Principles of remote attestation," *International Journal of Information Security*, vol. 10, no. 2, pp. 63–81, June 2011.
11. J. Ramsdell, P. D. Rowe, P. Alexander, S. Helble, P. Loscocco, J. A. Pendergrass, and A. Petz, "Orchestrating layered attestations," in *Principles of Security and Trust (POST'19)*, Prague, Czech Republic, April 8-11 2019.
12. H. Thiagarajan, J. Hatcliff, and Robby, "Awas: AADL information flow and error propagation analysis framework," *Innovations in Systems and Software Engineering (ISSE)*, 2021.
13. J. Hatcliff, J. Belt, Robby, and T. Carpenter, "HAMR: An

AADL multi-platform code generation toolset,” in *10th International Symposium on Leveraging Applications of Formal Methods, Verification and Validation (ISoLA)*, ser. LNCS, vol. 13036, 2021, pp. 274–295.

- 14. Robby and J. Hatcliff, “Slang: The Sireum programming language,” in *10th International Symposium on Leveraging Applications of Formal Methods, Verification and Validation (ISoLA)*, ser. LNCS, vol. 13036, 2021, pp. 253–273.
- 15. G. Klein, J. Andronick, K. Elphinstone, T. C. Murray, T. Sewell, R. Kolanski, and G. Heiser, “Comprehensive formal verification of an OS microkernel,” *ACM Trans. Comput. Syst.*, vol. 32, no. 1, pp. 2:1–2:70, 2014. [Online]. Available: <https://doi.org/10.1145/2560537>
- 16. G. Klein, J. Andronick, M. Fernandez, I. Kuz, T. C. Murray, and G. Heiser, “Formally verified software in the real world,” *Commun. ACM*, vol. 61, no. 10, pp. 68–77, 2018. [Online]. Available: <https://doi.org/10.1145/3230627>
- 17. I. Amundson and D. Cofer, “Resolute assurance arguments for cyber assured systems engineering,” in *Design Automation for Cyber-Physical Systems and Internet of Things (DESTION 2021)*, May 2021.

Darren Cofer is a Fellow at Collins Aerospace. His research interests include formal methods and tools for verification and certification of high-integrity systems. He received his PhD in Electrical and Computer Engineering from the University of Texas at Austin and is a senior member of the IEEE. Contact him at cofer@ieee.org.

Isaac Amundson is a Senior Principal Research Engineer at Collins Aerospace. His research focuses on modeling, formal verification and assurance, and he has led the design and implementation of multiple MBSE toolchains for providing inherent dependability in high-assurance systems. He received his PhD in Computer Science from Vanderbilt University.

Junaid Babar (PhD, Iowa State University) is a research engineer at Collins Aerospace. His research interests include formal verification via model checking, temporal logics and decision diagrams.

David Hardin is a research engineer at Collins Aerospace, and has made contributions in the areas of formal methods, high-assurance computer architecture, and real-time embedded Java. He is editor of the book *Design and Verification of Microprocessor Systems for High-Assurance Applications*. He received a PhD in Electrical and Computer Engineering

from Kansas State University.

Konrad Slind (PhD, TU Munich) is an industrial logician at Collins Aerospace. His research interests include design and implementation of higher order logic theorem provers, verified compilation, and self-describing data formats.

Perry Alexander is the AT&T Foundation Distinguished Professor and Director of The Information and Telecommunication Technology Center at The University of Kansas. His research interests include trusted computing, remote attestation, formal verification and formal synthesis. He received his PhD from The University of Kansas.

John Hatcliff is a University Distinguished Professor and Lucas-Rathbone Professor of Engineering in the Computer Science Department at Kansas State University. His research interests include software architecture, formal verification, and interoperability in the avionics, automotive, and medical domains. He received his PhD in Computer Science from Kansas State University.

Robby is a Professor of Computer Science at Kansas State University working in the area of formal methods, software engineering, and programming languages. He received a NASA Turning Goals Into Reality award in 2003, a NSF CAREER award in 2007, an ICSE 2000 Most Influential Paper award in 2010, and an ACM SIGSOFT Impact award in 2010.

Gerwin Klein is Chief Scientist at Proofcraft and Conjoint Professor at UNSW Sydney. His research interests include software verification, semantics of programming languages, and proof engineering. He is the architect of the seL4 microkernel verification which received the SIGOPS Hall of Fame Award. He received his PhD in Computer Science from TU Munich.

Corey Lewis is a Senior Proof Engineer at UNSW Sydney. He has been involved in all stages of the verification of the seL4 microkernel.

Eric Mercer is an Associate Professor at Brigham Young University. His research interests include software verification, model checking, static analysis, and programming languages. He received his PhD in Electrical Engineering from the University of Utah.

John Shackleton is a Senior Principal Research Scientist at Adventium Labs, focused on real-time

Department Head

embedded systems, cybersecurity, and model-based
system engineering and analysis.